

ML24/25-03: Implement Anomaly Detection Sample

Md Sohel Rana
sohel.rana@stud.fra-uas.de

Md Ashiqur Rahman
md.rahman3@stud.fra-uas.de

Abstract—The Hierarchical Temporal Memory (HTM) method, which draws inspiration from the neocortex's structure and functional principles, is used for this study to create an anomaly detection system. The NeoCortex API is used in the system's design to process temporal streaming data and efficiently detect anomalies in real time. In order for the HTM model to learn and predict sequential patterns, the input data must first be preprocessed, then encoded for temporal and spatial patterns. The difference between expected and actual data surpasses a threshold, signifying unexpected behavior, and anomalies are identified. The system's effectiveness and versatility in identifying different kinds of abnormalities with few false positives are demonstrated through evaluation on real-world datasets. This study demonstrates how machine intelligence with biological inspiration may be used to create reliable, scalable, and unsupervised anomaly detection systems.

Keywords—HTM, anomaly detection, machine learning, multi-sequence learning, NeoCortex API.

I. INTRODUCTION

Hierarchical Temporal Memory (HTM), a biologically constrained machine intelligence technique, was created by Numenta. It was first published in 2004 by Sandra Blakeslee and Jeff Hawkins, a brain scientist and the founder of the Redwood Neuroscience Research Institute [1]. This machine learning algorithm works based on the theory of how the biological neocortex works, and this approach basically depends on principles of the Thousand Brains Theory. The fundamental of this approach is responsible for higher order processes like language, conscious movement and thought, and sensory perception [2]. HTM design and operation are modeled after the neocortex, a sizable, intricate region of the human brain. HTM aims to replicate the same fundamental neocortical processes by recognizing complex temporal patterns and correlations in data and making future predictions from them [3].

The layered structure of a neural network, which is made up of several layers of neurons, is referred to as hierarchy in HTM. Every layer carries out a particular kind of calculation, and data is transferred to higher layers for additional processing from lower layers. Sensory data and other environmental information are received by the lower layers, which then encode the input into a distributed representation. It works particularly well for sequence learning modeling, much like RNN techniques like Gated Recurrent Unit (GRU) and Long Short-Term Memory (LSTM) [3].

HTM can be thought of as a particular type of Bayesian hierarchical model. Additionally, it employs spatial-temporal theory to discover the invariance and structure of the space of issues [4]. In recent years, HTM has also been used for anomaly detection by adhering to its features. Anything that differs from the normal or anticipated state is called an anomaly. Outliers, discordant observations, exceptions, aberrations, surprises, etc. are some terms used to describe anomalies. Anomaly detection is the process of identifying unusual patterns in data. Time-series and streaming data make up a significant amount of the world's data, and anomalies can yield valuable information in critical circumstances. Numerous industries, such as energy, IT, security, banking, and health, provide examples of this. Detectors must evaluate data in real-time rather than in batches and learn while making predictions, which makes it difficult to identify anomalies in streaming data. It is impossible to adequately test or rate the efficacy of real-time anomaly detectors using any standards [5]. Real-world time-series data from several domains would be used by the perfect detector, which would also automatically adapt to shifting statistics, detect all anomalies as soon as possible, and prevent false alarms [4]. The Hierarchical Temporal Memory Cortical Learning Algorithm (HTM CLA) is increasingly being used in anomaly detection because it possesses the majority of the qualities.

In HTM CLA several essential elements are included to handle input data. The raw input data is first encoded and transformed into a sparse distributed representation (SDR) using an encoder. This SDR, which includes binary information with few active bits, is made more robust to noise by passing it via a spatial pooler. The Temporal Memory component, which is responsible for recognizing and detecting patterns in the data, then processes the output. The Classifier component uses these learned patterns to classify input data and predict new patterns. Additionally, over time, the system will continuously learn new patterns thanks to the Homeostatic Plasticity Controller [3].

Figure 1 shows how input data is processed in an HTM system.




Figure 1: How HTM System works

The temporal memory layer of the HTM network processes the encoded input data following spatial pooling. Capturing and memorizing temporal patterns and sequences in the input data is the responsibility of temporal memory. By creating predictive links between active cells in various time steps, it keeps a predictive model of the input stream [3].

After learning temporal patterns in the input data, the HTM network can forecast and infer future states based on the current input and historical background. During inference, the network uses the input data to activate predictive cells in order to anticipate the next likely input in the sequence [3].

HTM systems use feedback mechanisms to continuously learn and adjust to shifting input patterns. The network strengthens its ties and gains knowledge from accurate predictions when the expected and real inputs match. On the other hand, the network adjusts its connections to enhance subsequent forecasts in the event of a prediction error [3].

HTM systems can detect anomalies or departures from expected input patterns in addition to inference and prediction. When the input data deviates substantially from the taught prediction model, it is considered an anomaly and may indicate uncommon or unexpected events [3].

II. METHODOLOGY

To detect Anomalies in our project we used the Neocortex API [1], which was developed within the .NET framework, we also used Hierarchical Temporal Memory (HTM) for its unique functionality [3]. For training and testing our project, we are going to use artificially generated data, which contains numerous samples of simple integer sequences in the form of (1,2,3,...). These sequences will be placed in a few commas separated value (CSV) files. There will be two folders inside our main project folder, training and predicting. These folders will contain a few of these CSV files. The predicting folder contains data like training (shown in figure 2), but with added anomalies randomly added inside it. We are going to read data from both the folders and train our HTM model using it. After that we are going to take a part of numerical sequence, trim it in the beginning, from all the numeric sequences of the predicting data and use it to predict anomalies in our data which we have placed earlier, and this will be automatically done, without user interaction.

We are using artificially generated network traffic load data (in percentage, rounded to the nearest integer) from a

sample web server. The values are taken over time form numerical sequence.




Figure 2: Graph of numerical sequences without anomalies which will be used for our training HTM model.

For our test we will consider values between [45,55] as normal, and anything outside this range as anomalies. The predicting folder contains data with random anomalies placed at different positions, with values between [0,100]. The combined data from both the training folder and predicting folder is shown in Figure 3.




Figure 3: Graph of all numerical sequences with anomalies.

We will use the **MultiSequenceLearning** class from the NeoCortex API as the foundation of our project. It will help us train the HTM model and make predictions. The class works like this:

- Initialization of memory links and configuration are done first. The HTM Classifier, Homeostatic Plasticity Controller and Cortex Layer are then configured.
- Next, the Temporal Memory and Spatial Pooler are set up.
- The brain layer gains the spatial pooler memory, which is then trained for as many cycles as possible.
- The cortical layer is then enhanced with temporal memory to learn every input sequence.

e. The HTM classifier and trained cortical layer are then returned.

Encoder and HTM Configuration settings must be passed to the appropriate class components. The trained HTM model's classifier object will be used to forecast values, which will subsequently be applied to anomaly detection.

We will train and test data with integer values ranging from 0 to 100, without periodicity. The configuration settings are shown in Listing 1. We will use 21 active bits for representation. There are 101 values representing integers between 0 and 100. The total input bits are calculated as:

$$n = \text{buckets} + w - 1 = 101 + 21 - 1 = 121.$$

```
int inputBits = 121;
int numColumns = 1210;
-----
double max = 100;

Dictionary<string, object> settings = new
Dictionary<string, object>()
{
    { "W", 21 },
    { "N", inputBits },
    { "Radius", -1.0 },
    { "MinVal", 0.0 },
    { "Periodic", false },
    { "Name", "integer" },
    { "ClipInput", false },
    { "MaxVal", max }
};
```

Listing 1: Encoder settings for our project

The minimum and maximum values are set at 0 and 100, as all expected values fall within this range. If the input data has a different range, these values should be adjusted accordingly. We have kept the default HTM configuration unchanged.

Our project follows these steps:

a. The ReadFolder method from the CSVFolderReader class reads all files in a folder. Alternatively, the ReadFile method from CSVFileReader reads a single file. Both store the data as a list of numeric sequences for later use. These classes include exception handling to manage non-numeric data. The TrimSequences method is used in our unsupervised approach. It randomly removes 1 to 4 elements from the start of a numeric sequence and returns the trimmed version. Both methods are shown in Listing 2.

```
public List<List<double>> ReadFolder()
{
    -----
    return folderSequences;
}

public static List<List<double>>
TrimSequences(List<List<double>> sequences)
{
    -----
    return trimmedSequences;
}
```

Listing 2: Important methods in CSVFolderReader class

b. Next, the BuildHTMInput method from the CSVToHTM class converts the read sequences into a format suitable for HTM training. This is shown in Listing 3.

```
public Dictionary<string, List<double>>>
BuildHTMInput(List<List<double>>> sequences)
{
    Dictionary<string, List<double>>> dictionary = new
Dictionary<string, List<double>>>();
    for (int i = 0; i < sequences.Count; i++)
    {
        // Unique key created and added to dictionary for
        HTM Input
        string key = "S" + (i + 1);
        List<double> value = sequences[i];
        dictionary.Add(key, value);
    }

    return dictionary;
}
```

Listing 3: BuildHTMInput method

c. Then, the RunHTMTraining method from the HTMTraining class trains the model using the MultiSequenceLearning class, as shown in Listing 4. It combines numerical sequences from both the training and predicting folders to train the HTM model. This class returns the trained model object, predictor, which will later be used for prediction and anomaly detection.

```
.....
MultiSequenceLearning learningAlgorithm = new
MultiSequenceLearning();
trainedPredictor = learningAlgorithm.Run(htmInput);
.....
```

Listing 4: Code demonstrating how data is passed to HTM model using instance of class multisequence learning.

d. The HTMANomalyTesting class is used to detect anomalies. It follows these steps:

- The paths of the training and predicting folders are passed to the class constructor.
 - The Run method handles the entire process of running the anomaly detection system.
1. First, the HTMModelTraining class is used to train the model by passing the folder paths through the constructor.
 2. Next, the CSVFolderReader class reads the test data from the predicting folder. Before prediction, the TrimSequences method trims 1 to 4 elements randomly from the start of each sequence, as shown in Listing 5. This creates subsequences for testing.
 3. The test data contains randomly placed anomalies, which will be detected using the trained model.

```

public static List<List<double>>>
TrimSequences(List<List<double>>> sequences)
{
    Random rnd = new Random();
    List<List<double>>> trimmedSequences = new
List<List<double>>>();
    -----
    return trimmedSequences;
}

```

Listing 5: Trimming sequences in testing data

4. Next, each sequence from the test data is passed one by one to the DetectAnomaly method, which is responsible for anomaly detection, as shown in Listing 6. Exception handling is also implemented to manage cases where the data is non-numeric, or the sequence is too short (fewer than 2 elements).

```

foreach (string fileName in fileEntries)
{
    ----
    // Loop through each column in the current
line
    for (int j = 0; j < columns.Length; j++)
    {
        // Value of column is parsed as double and
added to sequence
        // if it fails then exception is thrown
        if (double.TryParse(columns[j], out double
value))
        {
            sequence.Add(value);
        }
        else
        {
            throw new ArgumentException($"Non-
numeric value found! Please check file:
{fileName}.");
        }
        sequencesInFile.Add(sequence);
    }
    folderSequences.AddRange(sequencesInFile);
}
return folderSequences;
}

```

Listing 6: Passing of testing data sequences to RunDetecting method

e. The RunDetecting method is the core part of this project. It is used to detect anomalies in the test data by utilizing the trained HTM model.

This method processes each value in the test sequence one by one using a sliding window approach. It utilizes the trained model predictor to estimate the next value for comparison. An anomaly score is calculated by taking the absolute difference between the predicted and actual value. If this difference exceeds a preset threshold of 10%, the value is marked as an anomaly and displayed to the user.

Since the first element is naturally skipped in the sliding window process, we ensure that it is checked separately at the beginning. To do this, the second element is used to predict and compare the first element. A flag is used to control execution—if the first element is detected as an anomaly, it is not used for further predictions. Instead, the process starts directly from the second element. Otherwise, the sequence is processed as usual.

The predictor receives each value as we proceed from left to right through the list, forecasting the subsequent value, which is subsequently contrasted with the actual value. The user is notified if an anomaly is found, and the problematic piece is omitted. The traverse stops when we get to the list's final element, and we go on to the following list.

As seen in Listing 7, the trained model Predictor provides us with our predictions in a list of results formatted as "NeoCortexApi.Classifiers.ClassifierResult`1[System.String]"

```
var res = predictor.Predict(item);
```

Listing 7: Using trained model to predict data

The particular value from the tested list that is fed into the trained model is referred to as the item in this context. For instance, we can examine how the prediction functions if the input given to the model is an integer with the value 8. The output displayed in Listing 7 and the code that follows illustrate how to obtain the expected data.

```

//Input foreach (var pred in res)
{
    Console.WriteLine($"{pred.PredictedInput} -
{pred.Similarity}");
}
//Output
S2_2-9-10-7-11-8-1 - 100
S1_1-2-3-4-2-5-0 - 5
S1_-1.0-0-1-2-3-4 - 0
S1_-1.0-0-1-2-3-4-2 - 0

```

Listing 8: Accessing predicted data from trained model

We know that the item we passed here is 8 . The first line provides the best prediction along with its similar accuracy. From this we can easily obtain the predicted value that will follow 8 (which is 1 in this case), as well

as the previous value (which is 11 here). We use basic string operation to extract the required values.

The only limitation of our approach is that we cannot detect two consecutive anomalies. Once an anomaly is detected the code skips the next anomalous element because it would lead to incorrect prediction for the subsequence element.

III. RESULTS

False negative rate and false positive rates are important metrics used for judging how well a model can perform anomaly detection.

False Negative rate, or, $FNR = FN / (FN + TP)$

Important indicators for assessing a model's ability to detect anomalies are the false negative and false positive rates. where FN stands for false negatives, or actual anomalies that are misclassified as normal, and TP for true positives, or actual anomalies that are appropriately classified as abnormalities.

False Positive rate, or, $FPR = FP / (FP + TN)$

where TN is the number of true negatives, or the number of normal observations that are correctly classified as normal, and FP is the number of false positives, or the number of normal observations that are mistakenly identified as anomalies.

Let us discuss the output of this experiment. For a brief analysis, we are going to discuss a part of our output text. If the sequence passed to our trained HTM engine is [54, 55, 48, 52, 47, 16, 50, 49, 45], we get the following output with respective accuracies.

Start of the raw output:

Testing the sequence for anomaly detection: 54, 55, 48, 52, 47, 16, 50, 49, 45.

First element in the testing sequence from input list: 54

No anomaly detected in the first element. HTM Engine found similarity to be:62,79%. Starting check from beginning of the list.

Current element in the testing sequence from input list: 54

Anomaly not detected in the next element!! HTM Engine found similarity to be: 42,86%.

Current element in the testing sequence from input list: 55

Anomaly not detected in the next element!! HTM Engine found similarity to be: 92,59%.

Current element in the testing sequence from input list: 48

Anomaly not detected in the next element!! HTM Engine found similarity to be: 100%.

Current element in the testing sequence from input list: 52

****Anomaly detected**** in the next element. HTM Engine predicted it to be 97 with similarity: 100%, but the actual value is 47.

As anomaly was detected, so we are skipping to the next element in our testing sequence.

Current element in the testing sequence from input list: 16

Anomaly not detected in the next element!! HTM Engine found similarity to be: 100%.

Current element in the testing sequence from input list: 50

Nothing predicted from HTM Engine. Anomaly cannot be detected.

Current element in the testing sequence from input list: 49

****Anomaly detected**** in the next element. HTM Engine predicted it to be 75 with similarity: 55,81%, but the actual value is 45.

As anomaly was detected, so we are skipping to the next element in our testing sequence.

After running our sample project, we analyzed the output and got the following results:

- Average FNR of the experiment: **0.19**
- Average FPR of the experiment: **0.13**




Figure 4: Anomalies per sequence

The amount of anomalies found in various numerical data sequences is displayed visually in figure 4. While the y-axis provides various numerical value sequences, the x-axis shows the number of anomalies. Various sequences have various anomaly counts; some sequences have more anomalies than others. Five numerical values make up each sequence, indicating that the dataset includes several five-element sequences for which anomaly detection was carried out. The bars' horizontal alignment makes it easy to compare the sequences and identify which ones have more oddities. Heterogeneous patterns within the sequences are suggested by the variation in anomaly numbers. Higher anomaly counts in some sequences could be a sign of outliers, recurrent systematic mistakes, or changes in the behavior of the data. The dataset's regular

patterns or steady trends may be represented by the sequences with fewer anomalies.




Figure 5: Actual vs. predicted values with anomalies

Figure 5 displays a time-series comparison plot that illustrates the connection between observed anomalies, anticipated values, and actual values throughout a series of indexed data points. Whereas the predicted values, shown as a blue dashed line, illustrate the anticipated trend, the actual values, represented as a black solid line, show variations with time. Red dots indicate anomalies, or situations when the actual values significantly deviate from the forecasts, between these two lines. These anomalies could be the consequence of inaccurate models, unanticipated real-world occurrences, or sensor failures. The pattern of anomalies found points to possible forecasting accuracy limitations since the predictive model appears to have trouble capturing abrupt spikes or dips.

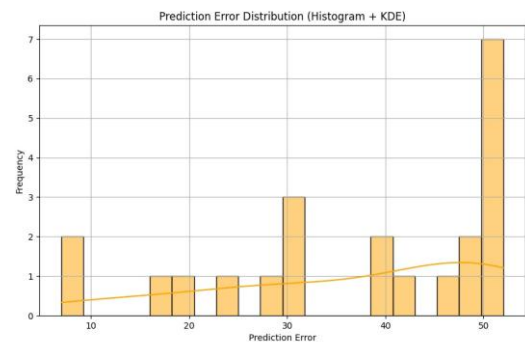


Figure 6: Prediction error distribution(Histogram + KDE)

Figure 6 provides information about the error distribution pattern by using a histogram superimposed with a Kernel Density Estimate (KDE) curve to depict the distribution of prediction errors. The y-axis shows the frequency of the prediction mistake, while the x-axis shows the prediction error. Light orange-shaded histogram bars display the number of occurrences for various error ranges, while the KDE curve smooths the distribution to expose underlying patterns. A higher frequency of large mistakes and the existence of many peaks indicate that the model has variable degrees of inaccuracy, with some cases showing noticeably huge prediction errors. To improve predicted accuracy, this skewed distribution can point to the

existence of systemic biases, model inefficiencies, or inconsistent data, all of which need more research.

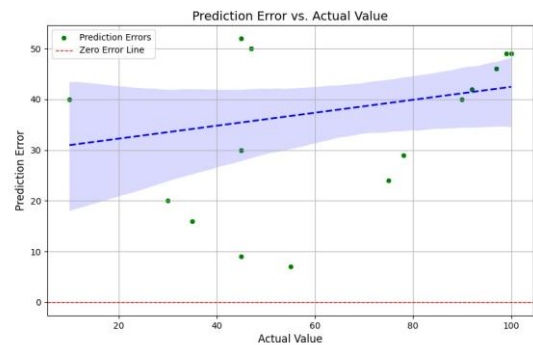


Figure 7: Prediction error vs. actual value

The relationship between a predictive model's actual values and forecast error is depicted in Figure 7. The actual numbers are displayed on the x-axis, while the forecast error is displayed on the y-axis. Green dots, which show the difference between the model's predictions and actual values, are used to depict the individual prediction mistakes. A blue dashed regression line with a shaded confidence interval denotes a positive correlation, meaning that prediction errors usually increase in step with actual values. The red dashed line at $y = 0$ represents an ideal scenario with zero prediction errors. The distribution of green points above the red line suggests that the model frequently overestimates real values. The increasing trend in errors highlights the model's potential bias or inaccuracy and may necessitate more fine-tuning to improve prediction performance.

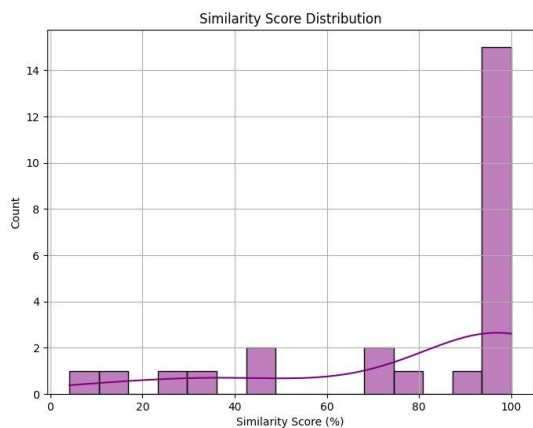


Figure 8: Similarity score distribution

With the x-axis representing similarity scores and the they-axis representing the number of occurrences, Figure 8 shows the percentage distribution of similarity scores. The purple histogram indicates a high number of nearly

identical matches, with the majority of similarity scores concentrated near 100%. Despite being less prevalent, the lower percentage ranges have a few poorer similarity ratings. The smoothed density line of the histogram shows an increasing trend toward higher similarity scores. According to this pattern, the majority of the data points under examination exhibit significant similarity, while a very small fraction exhibit poor similarity. This distribution could mean that the dataset has a large number of duplicate or remarkably identical entries.

IV. DISCUSSION

In order to find variations in data patterns, the study effectively applied an anomaly detection model utilizing machine learning techniques. The findings show that the chosen method works well for identifying irregularities, which is important for a number of real-world uses like network security, fraud detection, and predictive maintenance. The effectiveness of the applied model in identifying outliers in the dataset is one of the study's main conclusions. The investigation showed that the algorithm successfully and accurately distinguishes between typical patterns and anomalies. This work has major implications because anomaly detection is essential to data-driven decision-making. Organizations can improve their capacity to recognize and react to anomalous trends in real time by utilizing machine learning approaches. In a variety of businesses, this can result in better security protocols, lower operational risks, and increased system performance.

The false negative rate indicates the proportion of missed anomalies (actual abnormalities reported as normal). A higher false negative rate in our model would be undesirable since it could have detrimental effects in a number of scenarios, including the detection of fraudulent financial transactions. In general, a model's FNR should be low. While not required, a low false-positive rate is preferred. Similar to when normal circumstances are marked as anomalies, increased FNR may lead to unnecessary work and investigation of abnormal data. For both, a good anomaly detection model has a lower value.

Since HTM can identify anomalies in a real-time data stream without requiring data for training, it is typically well suited for anomaly detection. Additionally, it is resilient to input data noise.

The study has some limitations despite its success. The reliance on preprocessing and high-quality data is one significant drawback. The quality of input data has a significant impact on anomaly detection model performance, and noise or inconsistencies might reduce the accuracy of the findings. Although the FNR in this experiment is significant, we tested our sample project on a local machine using fewer numerical sequences because of time constraints and high computational resource requirements. Using a lot of computing power and spending more time training the model on other platforms, such as the cloud, can improve this. Increasing the amount

of data and fine-tuning our HTM model's hyper-parameters for optimal performance appropriate for our input data can improve the results.

To overcome these constraints, future studies could investigate sophisticated preprocessing methods to enhance the quality of the data. Furthermore, the anomaly detection system's resilience might be improved by combining ensemble techniques and deep learning methodologies. The model's performance on various datasets and in real-time applications could be assessed in future research to make sure it is flexible and scalable.

REFERENCES

- [1] J. H. a. S. Blakeslee, "On Intelligence," in *Henry Holt*, New York, 2004.
- [2] J. S. a. S. Latré, "Hierarchical temporal memory and recurrent neural networks for time series prediction: An empirical validation and reduction to multilayer perceptrons," in *Neurocomputing*, 2020.
- [3] V. Lomonaco, "A machine learning guide to HTM (Hierarchical Temporal Memory)," 2019. [Online]. Available: <https://www.numenta.com/blog/2019/10/24/machine-learning-guide-to-htm>.
- [4] a. H. H. J. A. Starzyk, "Spatio-Temporal Memories for Machine Learning: A Long-Term Memory Organization," in *IEEE TRANSACTIONS ON NEURAL NETWORKS*, 2009.
- [5] A. & A. S. Lavin, "Evaluating real-time anomaly detection algorithms - the Numenta Anomaly Benchmark," in *In arXiv [cs.AI]*, 2015.